

PRODUCT REQUIREMENTS DOCUMENT

AI-Powered After-Call Automation Platform

Prepared for Kalam CX

Document Owner: Product Management | Audience: Business Stakeholders & Executive Sponsors
Status: Draft for Review | Version 1.0 | Date: July 26, 2026

1. Product Overview

The AI-Powered After-Call Automation Platform is a new capability for Kalam CX designed to automatically generate structured, accurate, and consistent after-call records once a customer support interaction has ended. The platform reads the conversation between the agent and the customer, understands what the customer needed, identifies how the customer felt during the call, and produces a complete, organized summary that is ready for an agent to review and approve before it becomes part of the permanent customer record.

Rather than replacing agents, the platform removes the most repetitive and time-consuming part of their job: manually typing up notes after every call. It does this while preserving accountability, because no record is ever saved to the customer relationship system without a human confirming that it is correct.

2. Business Problem

Kalam CX agents currently spend a significant portion of every shift manually documenting completed calls. This after-call work is slow, mentally taxing after a long queue of conversations, and highly inconsistent from one agent to another. The consequences of this manual process include:

- Extended after-call handling time, which reduces the number of customers each agent can serve per day.
- Inconsistent record quality, with some agents writing detailed notes and others writing minimal or incomplete entries.
- Missed or inaccurate categorization of the reason for a call, the urgency of the issue, and the outcome reached.
- Delayed visibility for supervisors and quality teams, who must rely on incomplete notes to understand what actually happened on a call.
- Higher operating cost, as after-call work adds directly to labor hours without adding customer-facing value.

These issues compound at scale. With thousands of calls handled daily, even small inefficiencies in after-call documentation translate into substantial lost productivity, inconsistent reporting, and reduced ability to identify trends that matter to the business.

3. Product Vision

To make after-call documentation effortless, accurate, and instantaneous, so that Kalam CX agents can focus entirely on serving customers while every interaction is captured in a consistent, trustworthy, and audit-ready record — always reviewed and approved by a human before it becomes final.

4. Objectives

1. Reduce the time agents spend on after-call documentation by automatically drafting the record for their review.
2. Improve the consistency and completeness of customer records across all agents and shifts.
3. Increase the accuracy of call categorization, priority assignment, and outcome tracking for reporting and operational planning.
4. Preserve human accountability by requiring explicit agent approval before any record is finalized.
5. Provide leadership with clear, real-time visibility into after-call operations through operational statistics.
6. Ensure that automatically generated records never contain invented or unverifiable information.

5. Target Users

The platform is designed to serve four primary groups within Kalam CX's customer experience organization:

- Customer Support Agents — who handle live calls and are currently responsible for after-call documentation.
- Team Leaders / Supervisors — who oversee agent performance and review flagged or escalated records.
- Quality Assurance (QA) Specialists — who audit call records for accuracy, compliance, and coaching purposes.
- Operations & Business Leadership — who rely on aggregated call data and statistics for planning and decision-making.

6. User Personas

Persona	Role & Goals	Key Frustration Today
Amina — Frontline Agent	Handles 40–60 calls per day. Wants to close each call quickly and move to the next customer without sacrificing accuracy.	Spends several minutes after every call typing notes, which delays her availability for the next customer.
Karim — Team Leader	Manages a team of 12 agents. Needs to quickly understand what happened on escalated or sensitive calls.	Agent notes vary widely in quality, making it hard to assess situations quickly or coach effectively.
Layla — QA Specialist	Audits a sample of calls each week to ensure documentation meets compliance and quality standards.	Incomplete or inconsistent notes make it difficult to verify whether policy was followed correctly.
Youssef — Operations Director	Tracks call volume, category trends, and resolution rates to guide staffing and process decisions.	Reporting is only as good as the underlying notes, and inconsistent data undermines confidence in dashboards.

7. User Journey

The typical journey an agent experiences with the platform, from call completion to record finalization, is as follows:

- The customer call concludes and the conversation transcript becomes available to the platform.
- The platform analyzes the transcript to understand customer intent, sentiment, and relevant details discussed during the call.
- The platform retrieves the customer's existing information to provide context for the record.
- The platform checks relevant documentation and policy guidance to ensure the summary aligns with how similar cases should be handled.
- The platform generates a structured after-call record, clearly marking any field it cannot confidently determine as "Not Stated" rather than guessing.
- The platform classifies the call by category, priority, disposition, and customer sentiment, and determines whether the record should be Auto Saved, Routed to Review, or Escalated.
- The agent (or, where routed, a supervisor) reviews the draft record, makes any necessary corrections, and approves it.
- Only after approval is the record finalized and saved to the customer relationship system.
- Supervisors and leadership can view operational statistics reflecting call volumes, categories, sentiment trends, and review outcomes.

8. Functional Requirements

8.1 Transcript Intake

- The platform must accept a completed call transcript as the starting point for after-call processing.
- The platform must confirm that a transcript has been fully received before beginning analysis.

8.2 Understanding the Conversation

- The platform must identify the customer's underlying intent — the reason the customer contacted support.
- The platform must identify the customer's sentiment during the call (for example, satisfied, neutral, frustrated, or upset).
- The platform must retrieve the relevant customer's existing information to enrich the after-call record with accurate context.
- The platform must reference relevant documentation and policy rules to ensure the record aligns with established guidance for similar cases.

8.3 Record Generation

- The platform must generate a structured after-call record containing all standard fields used by Kalam CX (such as reason for contact, actions taken, outcome, and follow-up needs).
- The platform must never invent, assume, or fabricate information that is not supported by the transcript, customer data, or documentation.
- Any field that cannot be confidently populated from available information must be explicitly marked as "Not Stated" rather than left blank or guessed.

8.4 Classification

- The platform must classify each call by category (the type of issue or request).
- The platform must classify each call by priority (the urgency of the issue).
- The platform must classify each call by disposition (how the call concluded, such as resolved, pending, or unresolved).
- The platform must classify the overall customer sentiment for the interaction.

8.5 Routing Decision

- The platform must determine, for every record, whether it should be Auto Saved, Routed to Review, or Escalated, based on defined business rules.

8.6 Human Review and Approval

- The platform must present every generated record to a human for confirmation before it is treated as final, in line with the Human-in-the-Loop Workflow described in Section 11.
- The platform must allow the reviewing human to edit any field before approval.

8.7 Operational Statistics

- The platform must display operational statistics covering call volumes, category and priority distribution, sentiment trends, and the proportion of records Auto Saved, Routed to Review, or Escalated.

9. Non-Functional Requirements

Category	Requirement
Accuracy	Generated records must faithfully reflect only what is present in the transcript, customer information, and documentation, with no fabricated content.
Timeliness	After-call records must be generated promptly after a transcript becomes available, so agents can review and move on without meaningful delay.
Consistency	Records must follow a uniform structure and classification approach across all agents, teams, and shifts.
Reliability	The platform must operate dependably across the high daily call volumes handled by Kalam CX.
Auditability	Every record must retain a clear trail showing what was auto-generated, what was edited by a human, and who approved it.
Data Privacy & Security	Customer information must be handled in accordance with Kalam CX's data protection standards and applicable regulations at all times.
Scalability	The platform must support growth in call volume without degradation in turnaround time or record quality.
Usability	The review and approval experience must be simple enough that agents and supervisors require minimal training.

10. Product Features

Feature	Description
Transcript Intake	Accepts the finished call transcript as the trigger for after-call processing.
Intent Understanding	Identifies the core reason the customer contacted support.
Sentiment Detection	Identifies how the customer felt throughout the call and at its conclusion.
Customer Information Retrieval	Pulls in relevant existing customer details to give the record proper context.
Documentation & Policy Reference	Checks relevant guidance to keep the summary aligned with how similar cases are expected to be handled.
Structured Record Generation	Produces a complete, organized after-call record in the standard Kalam CX format.
"Not Stated" Safeguard	Marks any field that cannot be confidently determined as "Not Stated" instead of guessing.
Category, Priority, Disposition & Sentiment Classification	Automatically tags each record with these four business-critical classifications.
Smart Routing Decision	Automatically determines whether a record should be Auto Saved, Routed to Review, or Escalated.

Feature	Description
Human Review & Approval Workspace	Gives agents and supervisors a simple interface to check, correct, and approve records.
Operational Statistics Dashboard	Displays real-time and historical statistics on call volumes, categories, sentiment, and routing outcomes.

11. Human-in-the-Loop Workflow

A central principle of this platform is that automation supports human judgment rather than replacing it. No after-call record is ever saved to the customer relationship system without human confirmation. The workflow operates as follows:

16. Draft Generation: The platform generates a complete draft record, including classifications and a recommended routing outcome.
17. Routing Decision: Based on defined business rules (see Section 12), the record is automatically directed down one of three paths:
 - Auto Save Path — Low-risk, high-confidence records may proceed to the agent for a fast final confirmation.
 - Route to Review Path — Records with uncertain classifications, incomplete information, or moderate risk are sent to the agent or supervisor for closer review.
 - Escalation Path — Records involving high-priority issues, strong negative sentiment, or policy-sensitive situations are sent directly to a supervisor for review.
18. Human Review: The assigned reviewer examines the draft, makes any necessary corrections, and confirms accuracy.
19. Approval & Finalization: Only once a human has approved the record is it saved as final in the customer relationship system.
20. Traceability: The system retains a record of the original automated draft, any human edits, and the identity of the approver, supporting quality review and audit needs.

Guiding Principle: Automation drafts. Humans decide. No record becomes final without explicit human approval.

12. Business Rules

21. Every completed call must result in an after-call record; no call may be left undocumented.
22. A record may never be saved to the customer relationship system without human approval, regardless of the routing outcome.
23. Any information the platform cannot confidently confirm must be marked "Not Stated"; the platform must never fabricate details to complete a field.
24. Calls classified with high priority, strongly negative sentiment, or policy-sensitive content must be Escalated to a supervisor rather than Auto Saved.
 1. Records with incomplete customer information, ambiguous intent, or conflicting details must be Routed to Review rather than Auto Saved.
 2. Only records that meet defined confidence and completeness thresholds are eligible for the Auto Save path, and even then require a final human confirmation.
 3. Category, priority, disposition, and sentiment classifications must be recorded for every call, regardless of routing path.
 4. All edits made by a human reviewer must be captured for audit and quality purposes.
 5. Operational statistics must be refreshed regularly to reflect current call activity and routing outcomes.

13. Success Metrics

Metric	Target Outcome
After-call handling time per call	Meaningful reduction compared to fully manual documentation.
Record completeness rate	High proportion of records with all required fields populated or explicitly marked "Not Stated".
Classification accuracy	High agreement between platform-assigned category, priority, disposition, and sentiment and human-confirmed values.
Auto Save rate	Steady, sustainable share of eligible records proceeding through the fast-confirmation path without compromising quality.
Escalation appropriateness	High proportion of escalated records confirmed by supervisors as correctly escalated.
Agent satisfaction	Improved agent-reported satisfaction with after-call workload.
Data integrity incidents	Zero confirmed cases of fabricated or unsupported information being saved as final.

14. Acceptance Criteria

- The platform successfully accepts a call transcript and produces a structured after-call record for every completed call.
- Generated records correctly identify customer intent and sentiment in line with human reviewer judgment during acceptance testing.
- Generated records correctly incorporate relevant customer information and applicable documentation guidance.
- Any field lacking sufficient supporting information is clearly marked "Not Stated," with no fabricated content present in any tested record.
- Every record includes a category, priority, disposition, and sentiment classification.
- Every record is assigned a routing outcome of Auto Save, Route to Review, or Escalate, consistent with the defined business rules.
- No record can be saved to the customer relationship system without a recorded human approval action.
- Operational statistics accurately reflect processed call volumes, classification distributions, and routing outcomes.

15. Assumptions

- A complete and legible call transcript will be available for each call shortly after the call concludes.
- Existing customer information is accessible and reasonably up to date at the time of record generation.
- Current documentation and policy guidance used for reference is maintained and kept current by relevant business teams.
- Agents and supervisors will receive appropriate onboarding on the review and approval workspace.
- The standard after-call record format currently used by Kalam CX will serve as the basis for the structured record.

16. Constraints

- The platform must operate within Kalam CX's existing data protection and confidentiality requirements.
- The platform depends on the quality and completeness of the call transcript; poor-quality transcripts may limit the accuracy of generated records.
- Human approval is a mandatory step for every record, which sets a floor on how quickly records can be finalized.
- Initial rollout may be limited to a defined set of call categories or teams before wider deployment.

17. Risks

Risk	Mitigation
Agents over-rely on automated drafts without careful review.	Design the review workspace to make edits easy and require active confirmation, not passive acceptance; monitor edit rates.
Inconsistent or missing documentation leads to lower-quality guidance references.	Establish an ongoing process to keep policy and documentation sources current and complete.
Sentiment or intent misclassification affects escalation accuracy.	Apply conservative routing thresholds so uncertain cases default to Review or Escalation rather than Auto Save.
Change resistance from agents concerned about job impact.	Communicate clearly that the platform assists rather than replaces agents, and involve agents in rollout feedback.
Customer data handling concerns.	Enforce strict adherence to existing data protection and access control standards throughout the workflow.

18. Out of Scope

- Automated handling of live, in-progress calls; the platform operates only after a call has concluded.
- Fully autonomous saving of records without any possibility of human review.
- Customer-facing features or direct customer interaction of any kind.
- Replacing existing customer relationship system functionality beyond after-call record creation.
- Real-time agent coaching or in-call guidance.

19. Future Enhancements

- Expanding coverage to additional call categories and business units beyond the initial rollout.

-
- Trend analysis that surfaces recurring customer issues to inform product and policy improvements.
 - Personalized coaching insights for agents based on patterns observed across their after-call records.
 - Proactive quality sampling that highlights records most valuable for QA review.
 - Deeper integration of operational statistics into broader business planning and workforce management processes.

— *End of Document* —